

Minimising a Noisy Expensive Function Using Active Learning



Ross Brown

Supervisor: Dr. O. Orhobor

Department of Chemical Engineering and Biotechnology
University of Cambridge

This dissertation is submitted for the degree of
Master of Engineering

Declaration

I hereby declare that except where specific reference is made to the work of others, the contents of this dissertation are original and have not been submitted in whole or in part for consideration for any other degree or qualification in this, or any other university. This dissertation is my own work and contains nothing which is the outcome of work done in collaboration with others, except as specified in the text and Acknowledgements. This dissertation contains fewer than 65,000 words including appendices, bibliography, footnotes, tables and equations and has fewer than 150 figures.

Ross Brown
September 2021

Abstract

Experimentation is a tedious and expensive process with many experimentations discarded as failures. Too often, the experimenter is focused too heavily on a local minimum that they believe holds all their answers. This narrow-minded approach leads to the loss of the global minimum: a far more fulfilling find. By using active learning, this paper aims to show that the global minima can be found with limited loss of resources and time. By using smoothing splines, the experimenter might not even test at the minimum in order to be confident on its location.

This paper focuses primarily on problems with 1 parameter. In doing so, it shows how for a noisy expensive function, the global minima can be found with fewer function calls than `fminbound` [2] and to a higher degree of certainty. Further work is encouraged within this area, particularly with the generalisation to n -dimensional problems.

Table of contents

List of figures	ix
1 Introduction	1
1.1 Problem Definition	1
1.2 Principles of Active Learning	1
2 Simple 1-Dimensional Problem	3
2.1 Outlining of Basic Principles	3
2.1.1 Algorithms	3
2.1.2 Comparison One	5
2.2 Deeper Analysis	7
2.2.1 Random Function	7
2.2.2 Comparison on Random Functions	7
3 Final Discussion	9
3.1 Further Work	9
3.2 Conclusion	9
References	11

List of figures

2.1	First Function	4
2.2	First Comparison	6
2.3	Rigorous Comparison	8

Chapter 1

Introduction

Finding the global minimum of a function within a set of boundaries is a problem of major import. From optimising a synthetic pathway in drug development, to minimising the error in a neural network, minimisation is vitally important to mathematics. Within the numerical field, the goal is usually two-fold: reduce the error, ε , to the true value *and* reduce the processing time. With these goals in mind, the majority of algorithms exploit the commonality of the cheapness of the target function. However, this is not always the case. Take as an example an experimentation of sand grain size, d , on the strength of concrete, τ . An underlying function of the form $\tau = f(d)$ exists, but each call to this function takes at least a day, and is labour and material expensive. The target of this paper is to explore how to minimise such a function with the fewest function calls.

1.1 Problem Definition

$$y = f(\mathbf{x}) \tag{1.1}$$

Given (1.1) where \mathbf{x} is a vector with $x_i \in [\alpha_i, \beta_i]$ and y is scalar, find the solution to $\operatorname{argmin}_{\mathbf{x}} [f(\mathbf{x})]$. The algorithm will be able to invoke $g(\mathbf{x})$ as shown in (1.2) where ε represents an unknown random error.

$$g(\mathbf{x}) = f(\mathbf{x}) + \varepsilon \tag{1.2}$$

1.2 Principles of Active Learning

Active learning involves the intelligent sampling with the intent of reducing the total number of labelling required. Uses of this can be seen in image recognition, where large sets of

images may exist but only a handful of these images have been fully labelled. Each additional image requires the employment of a human to interpret the image. In order to reduce the amount of labelling, [1] suggests there are three methods for reducing the number of required labels.

- Highest Uncertainty
- Competing Hypothesis
- Predicted Model Change

Here, the greatest uncertainty value implies knowledge to the probability of the output. Competing hypothesis theories involves several different models, where samples are chosen to test these hypotheses against each other. Finally, the predicted model change takes samples which are expected to have the largest impact on the modelling.

Chapter 2

Simple 1-Dimensional Problem

2.1 Outlining of Basic Principles

By constraining \mathbf{x} to one dimension allows for the problem to be simplified. Suppose $f(x) = \sin(x) + 0.05x^2$ with $x \in [-10, 10]$ as shown in Figure 2.1. In this range, there are multiple minima with only one global minimum. The task here is to successfully locate the minima situated at -1.428. Three methods will be used here: fminbound [2], greatest uncertainty active learning, and problem specific active learning (explained in Section 2.1.1). ϵ will be chosen to be independent of x and y and fit a normal distribution such that $\epsilon \sim N(0, 0.2^2)$.

2.1.1 Algorithms

fminbound

fminbound is a function included in the scipy optimisation library [2]. It uses Brent's method allowing it to be quick in situations where labelling is quick and error is low.

Greatest Uncertainty Active Learning

Without a basis for implementing a probability model on curve fitting, regions with maximum scarcity of samples will be chosen. In practice, this means for $n + 1$ samples within the range $[\alpha, \beta]$, the sample $s_i, i \in \mathbb{Z}, i \in [0, n]$ will fall according to (2.1). A smoothing spline fit will be used for interpolation identically to problem specific active learning in Section 2.1.1.

$$s_i = \alpha + i \frac{\beta - \alpha}{n} \quad (2.1)$$




Fig. 2.1 $y = \sin(x) + 0.005x^2$ with $x \in [-10, 10]$

Problem Specific Active Learning

This methodology has two underlying core principles: sparse areas reveal the most information and minimal areas reveal information to the location of minima. Combining these allows for better decision making in selecting the next sample. A smoothing spline is used to interpolate providing $e(x)$ in (2.2). Functions $h(x)$ and $p(x)$ (see (2.2) and (2.3)) allow for the minimal areas and most sparse areas to be found respectively. The next sample is chosen at $\operatorname{argmax}_x [h(x)p(x)]$.

$$h(x) = \frac{-e(x) + \max[e(x)]}{\max[e(x)] - \min[e(x)]} + 0.01 \quad (2.2)$$

$$p(s_i \leq x \leq s_{i+1}) = \min[x - s_i, s_{i+1} - x] \quad (2.3)$$

The smoothness attached to $e(x)$ is defined differently on each step. Here, all steps have a smoothness factor of $0.01n$ until the final estimation is made where the smoothing coefficient becomes $\frac{n}{30}$. The reasoning behind these steps comes down to tolerance. At low n , there simply are not enough data for ε to be noteworthy. Indeed, more courageous guesses are wanted while far from the maximum number of experimentations. There are two ways to achieve this: weight $p(x)$ higher or allow higher curvature on $e(x)$. The latter is focused upon in this paper. As more data appears, this need reduces and the smoothness can be increased. For the final iteration, a smoother curve is desired so scarce data and high uncertainty does not negatively affect the result.

2.1.2 Comparison on $\sin(x) + 0.005x^2$

Each method discussed in Section 2.1.1 was executed 50 times for each sample size between 2 and 25. Figure 2.2 shows the mean error, standard deviation, and absolute mean error of each method. When greater than nine samples were used, `fminbound` [2] fails to present itself as a competitive candidate. For sample sizes beneath this, the problem specific approach appears worse. This is due to the discovery of the local minima at 4.27. If this were indeed the global maxima, `fminsearch` may have failed to find the global minima in all cases. With respect to the two active learning methods, the problem specific method did slightly better, although this could be attributed to chance based on the curve chosen.

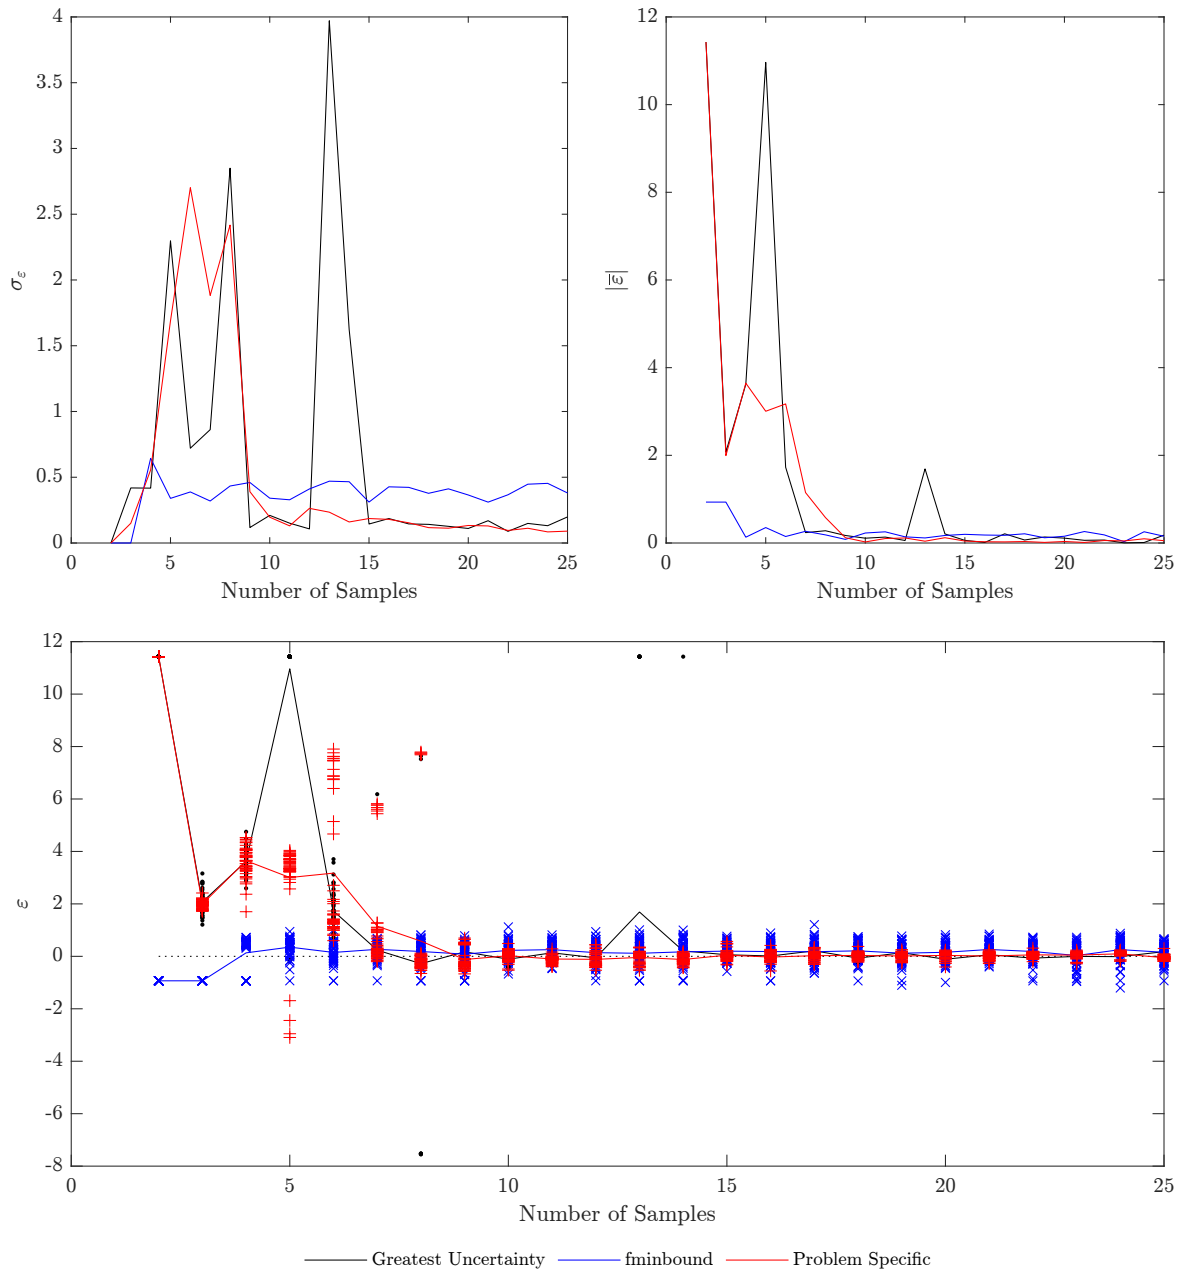


Fig. 2.2 Comparison of the three methods discussed. The top left shows the standard deviation of each method, the top right shows the absolute mean error, and the bottom shows the true error of each sample with a mean of the error to guide the eye.

2.2 Deeper Analysis

2.2.1 Random Function

To fully test these methods, random functions are used. For this, several basic assumptions are made:

- The function is continuous.
- $x \in [-10, 10]$.
- $f(-10) = 0$.
- $f'(x) \sim N(0, 0.5^2)$.
- No additional error is attached (very noisy anyway).

A simple script for this is in Listing 2.1, which has been designed to allow for easy production of a curve that is differentiable to the n^{th} degree.

Listing 2.1 Funcion to generate a random continuous function.

```
def randFunc():
    lims = [-10, 10]
    derivatives = 1
    x = np.linspace(*lims, 5000)
    f = np.zeros([derivatives + 1, len(x)])
    f[-1, :] = np.random.normal(0, 0.5, np.shape(x))

    for i in range(1, derivatives):
        f[i, 0] = np.random.uniform(-1, 1)

    for i in range(1, len(x)):
        for j in range(derivatives):
            f[j, i] = f[j, i - 1] + f[j + 1, i - 1] * (x[i] - x[i - 1])

    return interp1d(x, f[0, :])
```

2.2.2 Comparison on Random Functions

As in Section 2.1.2, a comparison between the three methods was carried out on samples sizes ranging from 2 to 25. These were tested on 50 different random functions, as generated by Listing 2.1. The problem specific algorithm surpassed the other methods in both metrics

measured. Further, both active learning methods routinely surpassed `fminbound`. This is likely due to `fminbound` finding local minima and converging upon these rather than exploring the remainder of the function.

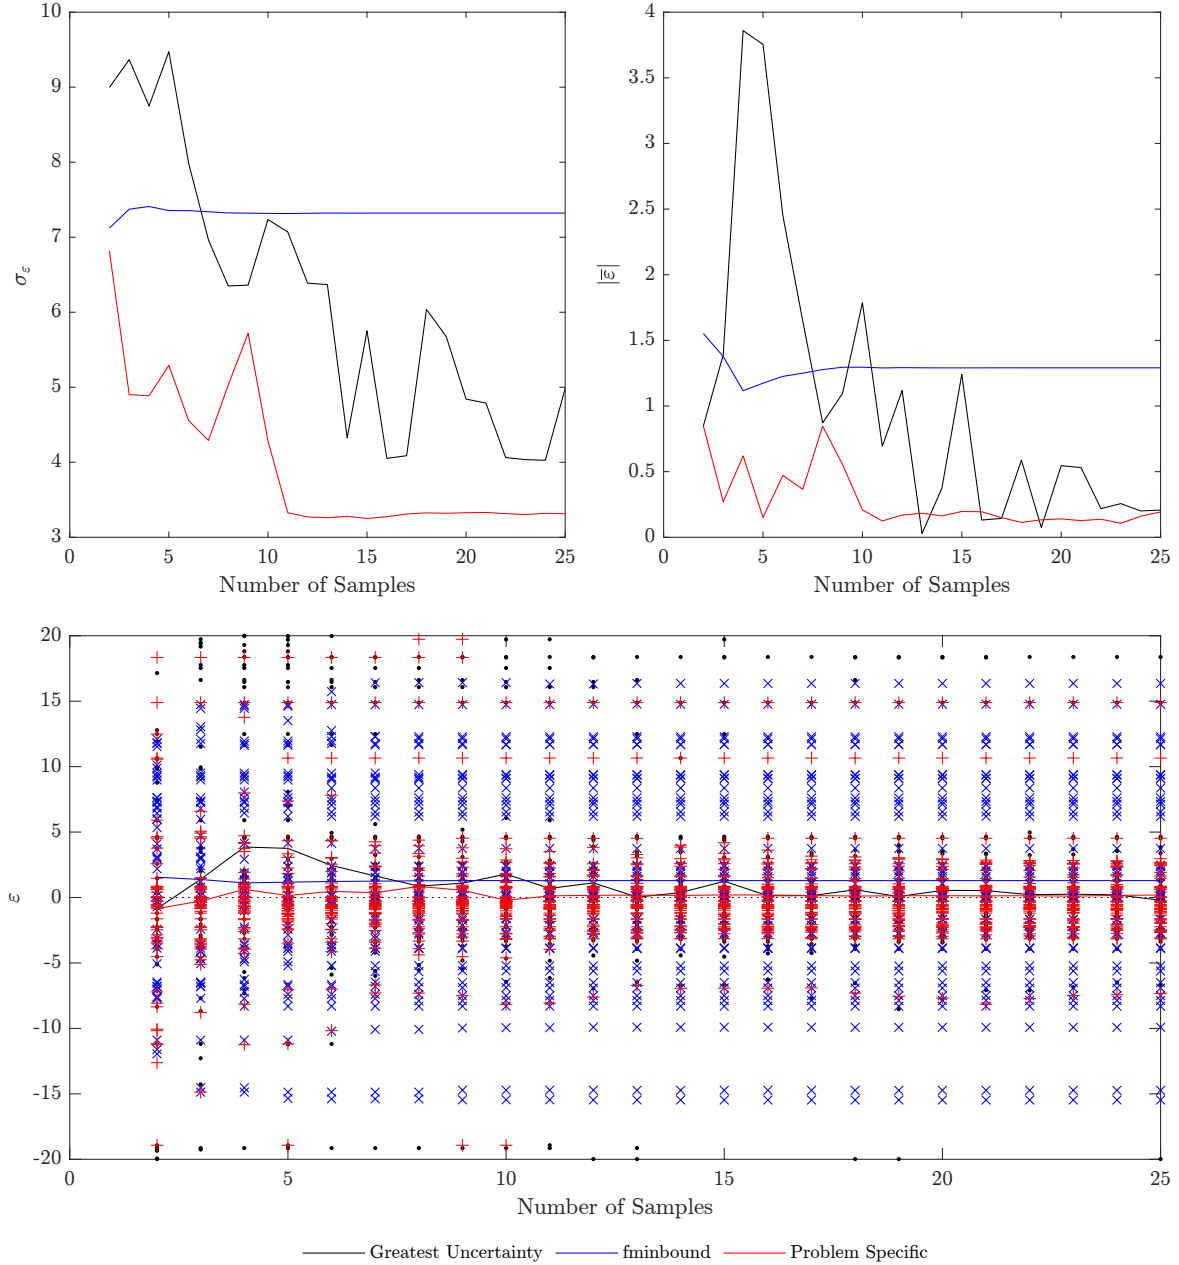


Fig. 2.3 Comparison of the three methods discussed against random functions. The top left shows the standard deviation of each method, the top right shows the absolute mean error, and the bottom shows the true error of each sample with a mean of the error to guide the eye.

Chapter 3

Final Discussion

3.1 Further Work

Allowing for n -dimensional minimisation follows routine procedures. Indeed, simply using `splprep` from the `scipy.optimize` package allows for this to be modified to accept n -dimensions. All other procedures remain the same. This was not evaluated within this report as random n -dimensional functions difficult to produce. Further work was taken to attempt to converge upon suitable smoothing factors as a function of sample size, although this proved to be unreliable.

Investigation into different non-parametric curve fitting may be made. It is expected that a combination of locally estimated scatterplot smoothing (LOESS) and smoothing splines may be used: LOESS for heavily sampled regions and smoothing splines for sparser areas. Alternatively, regression bins may be used under the assumption that the minima can be approximated to quadratic. Upon estimation of the width of the minima, a suitable bin size may be determined.

Care is needed when choosing values minima with low data representation. This is the cause for many of the extreme deviations seen within the problem specific active learning methodology. To combat this, a coupling of estimated minima and certainty could be tailored.

3.2 Conclusion

Within this report, it has repeatedly been shown that active learning plays a role in optimisation in situations where noisy data and expensive function calls are present. It also demonstrates the information that is available during experimentation, which may be undervalued. By exploiting this information, time and resources may be saved.

Further efforts are needed to improve this methodology. Importantly, using this on n -dimensional problems since virtually all experiments have more than one parameter. Secondly, different fits should be compared, potentially even leading to competing hypothesis active learning.

References

- [1] Settles, B. (2009). Active learning literature survey. pages 1–47.
- [2] Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D., Burovski, E., Peterson, P., Weckesser, W., Bright, J., van der Walt, S. J., Brett, M., Wilson, J., Millman, K. J., Mayorov, N., Nelson, A. R. J., Jones, E., Kern, R., Larson, E., Carey, C. J., Polat, İ., Feng, Y., Moore, E. W., VanderPlas, J., Laxalde, D., Perktold, J., Cimrman, R., Henriksen, I., Quintero, E. A., Harris, C. R., Archibald, A. M., Ribeiro, A. H., Pedregosa, F., van Mulbregt, P., and SciPy 1.0 Contributors (2020). SciPy 1.0: Fundamental Algorithms for Scientific Computing in Python. *Nature Methods*, 17:261–272.